


# Apache Airflow and Amazon Redshift

Applications to Machine Learning and Analytics

Stephanie Kirmer

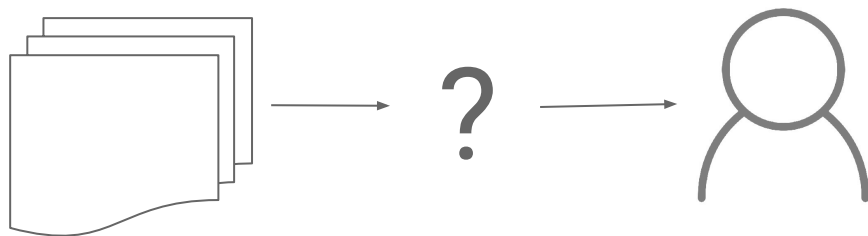
[www.stephaniekirmer.com](http://www.stephaniekirmer.com)

[https://github.com/skirmer/airflow\\_plus\\_redshift](https://github.com/skirmer/airflow_plus_redshift)

 @data\_stephanie

# Learning Goals

- ❖ What are Airflow and Redshift?
- ❖ How can they be used to build and manage a data warehouse?
- ❖ How can machine learning workflows employ them?
- ❖ What are the risks and pitfalls to avoid?



# Toolkit

Data storage solution:



- Installation instructions:  
<https://docs.aws.amazon.com/redshift/latest/gsg/getting-started.html>

Job scheduler:



- Installation instructions:  
<https://airflow.apache.org/docs/stable/start.html>

Optional, nice to have: a helpful Ops colleague

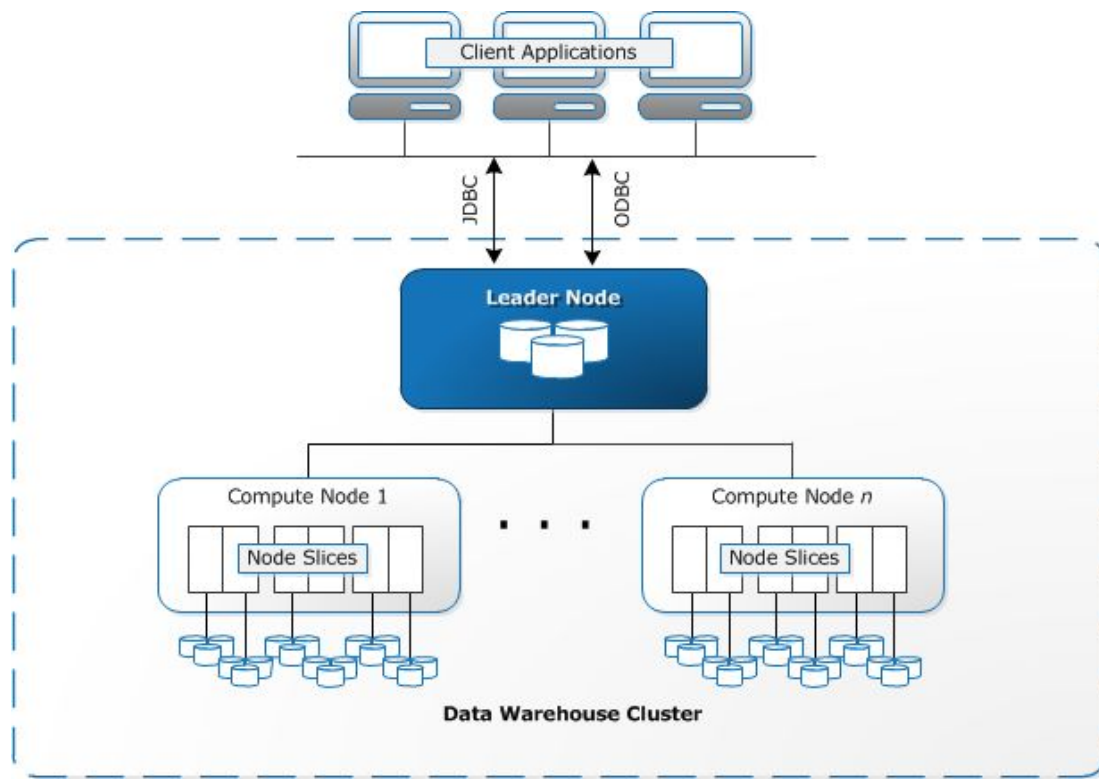
# Amazon Redshift



# Redshift is a data storage product from Amazon.

- ❖ Common ancestor with PostgreSQL - branched off from PostgreSQL 8.0.2
- ❖ Columnar storage, which allows powerful compression
- ❖ Gives lots of opportunities for parallelization because of distributed nature, but data manipulation clashes still can happen
- ❖ AWS tools allow you to programmatically create and destroy instances/clusters

# Architecture Overview



**Leader** receives queries and develops plans for execution.

**Compute Nodes** receive instruction from Leader and hand off to Slices to run.

Data is split into **Slices**, divided according to your chosen distribution keys.

[https://docs.aws.amazon.com/redshift/latest/dg/c\\_high\\_level\\_system\\_architecture.html](https://docs.aws.amazon.com/redshift/latest/dg/c_high_level_system_architecture.html)

Distributed data storage allows for rapid retrieval if organized well.

# Take these steps to make Redshift work for you.

## Understand the underlying framework

- ❖ Relational and distributed
- ❖ Distribution key helps Redshift choose what data to store together in nodes.
- ❖ Data compression- set encodings that make sense for your data (`ANALYZE COMPRESSION`).

## Plan ahead and be strategic

- ❖ It's difficult to change architecture midstream.
- ❖ Sort key: what do you usually filter/sort by? This might be it.
- ❖ Redshift wants to help you- use `EXPLAIN` to test how beefy your query is going to be.

# A Moment on Schemas

Redshift will allow you to build unkeyed, unlinked tables willy-nilly - *I advise against this!*

Employing a systematic, planned data architecture schema will save you pain and stress in the future.

Redshift does not enforce foreign or primary key restrictions the way other SQL types may - you have to maintain your own discipline.

Using table keys can make choosing highly performant sort keys and dist keys easier too!



# Tell Redshift what to expect.

```
CREATE TABLE IF NOT EXISTS public.customer(  
    customer_key integer not null IDENTITY(1,1) encode zstd sortkey,  
    customer_uuid char(36) encode zstd,  
    customer_home_base char(20) encode zstd,  
    customer_home_base_key integer encode zstd,  
    primary key(customer_key));
```

When creating long term persistent tables, set your schema in advance.

# Protect yourself when inserting data.

```
COPY public.customer(customer_uuid, customer_home_base, customer_home_base_key)
  from 's3://s3_bucket/s3_key_customer.csv'
  iam_role '{self.iam}'
  region 'us-east-1'
  removequotes
  NULL 'None'
  emptyasnull
  blanksasnull
  delimiter ','
;
```

When INSERT-ing or COPY-ing, explicitly state the column names – data coming from elsewhere may be ordered wrong.

# Mistakes I've Made With Redshift

## **Not optimizing queries to match keys**

- ❖ No query should run for 66 hours.

## **Making a table for every occasion**

- ❖ Think about how the table fits with your architecture. You have an architecture, right?
- ❖ Don't gunk up your data warehouse with tables that are not documented or reusable. Clean house sometimes!

## **Making tables without keys or clear relationships**

- ❖ Optimize, optimize, optimize - you will have more data, and more users, eventually

## **Making tables without compression**

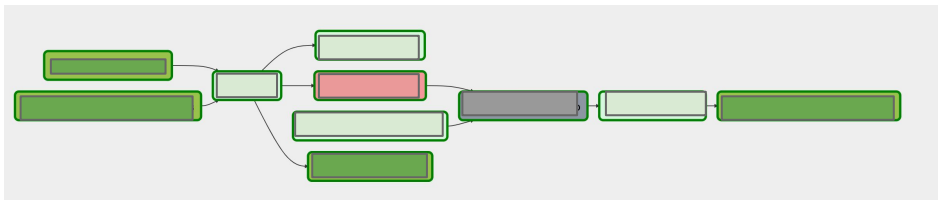
- ❖ This is one of the big advantages of Redshift, so why not use it?

# Apache Airflow

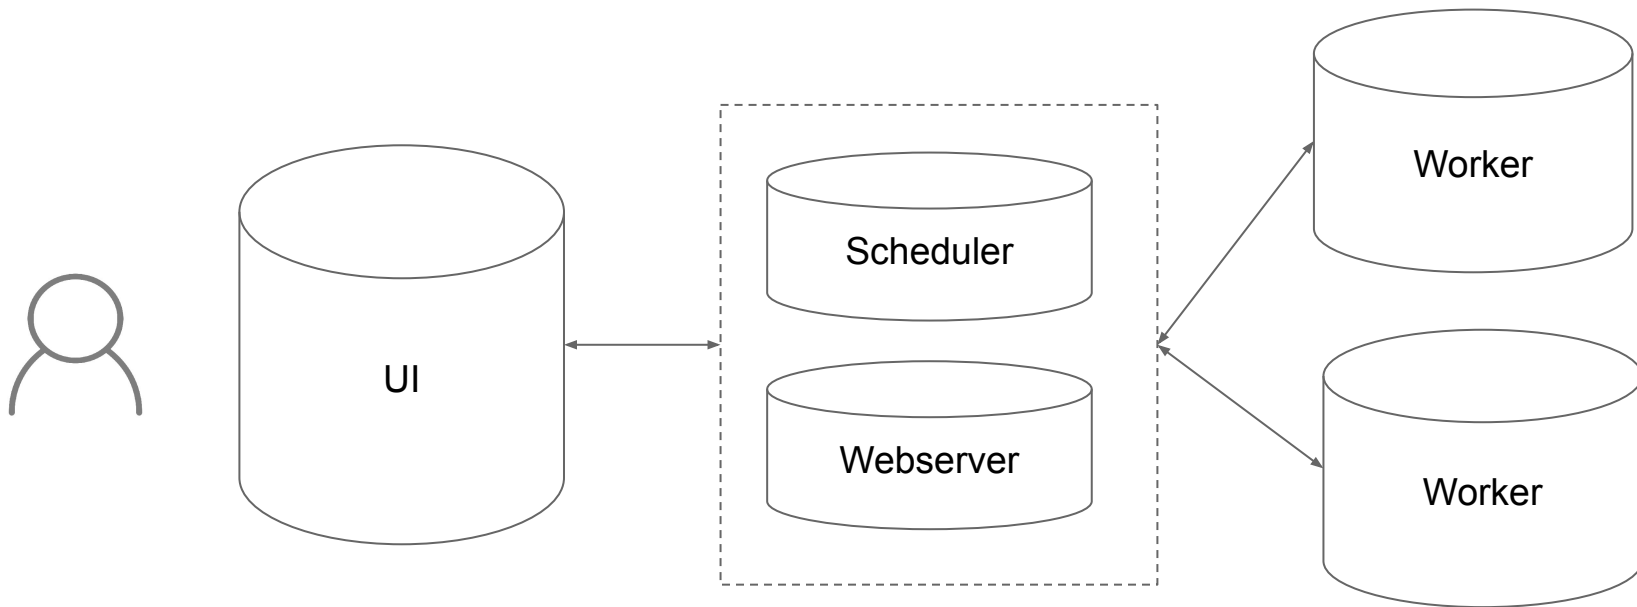


# Airflow is an open source task management tool.

- ❖ Tool for programmatically building, scheduling, and maintaining jobs
- ❖ Python based, uses Directed Acyclic Graphs (DAGs)
- ❖ Allows testing, versioning, and monitoring
- ❖ Open source and free - you can contribute!



# Airflow is structured in clusters.



You might need a friendly Ops person to help with installation. Note: I have omitted some elements of metadata management and logging here for simplicity's sake.

# Meet the Directed Acyclic Graph.

Airflow wants to know...	Write in your DAG definition file...
<i>What tools/resources do you need?</i>	Calls of python libraries, airflow modules and operators, etc
<i>What information will all your tasks want to know?</i>	Dictionary of default arguments- retry rules, failure notifications, etc
<i>What is this DAG all about?</i>	Call DAG function with arguments of DAG characteristics- name, run schedule, etc
<i>What is each task you want to perform?</i>	Call each operator, and define each task with the arguments
<i>What order should things happen?</i>	Define dependencies (task >> next_task)

For further reading: <https://airflow.apache.org/docs/stable/tutorial.html#it-s-a-dag-definition-file>

# Meet the Directed Acyclic Graph.

Airflow wants to know...	Write in your DAG definition file...
<i>What tools/resources do you need?</i>	<b>Calls of python libraries, airflow modules and operators, etc</b>
<i>What information will all your tasks want to know?</i>	Dictionary of default arguments- retry rules, failure notifications, etc
<i>What is this DAG all about?</i>	Call DAG function with arguments of DAG characteristics- name, run schedule, etc
<i>What is each task you want to perform?</i>	Call each operator, and define each task with the arguments
<i>What order should things happen?</i>	Define dependencies (task >> next_task)

For further reading: <https://airflow.apache.org/docs/stable/tutorial.html#it-s-a-dag-definition-file>



Airflow standards

```
from airflow import DAG
```

```
from airflow.contrib.operators.aws_athena_operator import AWSAthenaOperator
```

Other python  
libraries

```
from datetime import datetime, timedelta
```

Custom written  
operators

```
from operators.trigger_error_alert import trigger_pagerduty_alert
```

```
from operators.trigger_error_alert import resolve_pagerduty_alert
```

```
from operators.redshift_runner import RedshiftOperator
```

Custom SQL  
saved in separate  
.py file

```
from sql import records_query as hdq
```

# Meet the Directed Acyclic Graph.

Airflow wants to know...	Write in your DAG definition file...
<i>What tools/resources do you need?</i>	Calls of python libraries, airflow modules and operators, etc
<b><i>What information will all your tasks want to know?</i></b>	<b>Dictionary of default arguments- retry rules, failure notifications, etc</b>
<i>What is this DAG all about?</i>	Call DAG function with arguments of DAG characteristics- name, run schedule, etc
<i>What is each task you want to perform?</i>	Call each operator, and define each task with the arguments
<i>What order should things happen?</i>	Define dependencies (task >> next_task)

For further reading: <https://airflow.apache.org/docs/stable/tutorial.html#it-s-a-dag-definition-file>

```
default_args = {  
    "owner": "airflow",  
    "start_date": datetime(2020, 4, 5),  
    "depends_on_past": False,  
    "email": ["stephanie@stephaniekirmer.com"],  
    "email_on_failure": True,  
    "email_on_retry": False,  
    "retries": 3,  
    "retry_delay": timedelta(minutes=1),  
    "on_failure_callback": partial(trigger_pagerduty_alert, environ),  
    "on_success_callback": partial(resolve_pagerduty_alert, environ),  
    "provide_context": True,  
}
```

# Meet the Directed Acyclic Graph.

Airflow wants to know...	Write in your DAG definition file...
<i>What tools/resources do you need?</i>	Calls of python libraries, airflow modules and operators, etc
<i>What information will all your tasks want to know?</i>	Dictionary of default arguments- retry rules, failure notifications, etc
<i>What is this DAG all about?</i>	<b>Call DAG function with arguments of DAG characteristics- name, run schedule, etc</b>
<i>What is each task you want to perform?</i>	Call each operator, and define each task with the arguments
<i>What order should things happen?</i>	Define dependencies (task >> next_task)

For further reading: <https://airflow.apache.org/docs/stable/tutorial.html#it-s-a-dag-definition-file>

```
dag = DAG(  
    dag_id="my_dag_name",  
    default_args=default_args,  
    schedule_interval="30 3 1 * *",  
    catchup=True  
)
```

Note: it accepts the default args here

Scheduling intervals are in CRON syntax: [crontab.guru](http://crontab.guru) is a big help

# Meet the Directed Acyclic Graph.

Airflow wants to know...	Write in your DAG definition file...
<i>What tools/resources do you need?</i>	Calls of python libraries, airflow modules and operators, etc
<i>What information will all your tasks want to know?</i>	Dictionary of default arguments- retry rules, failure notifications, etc
<i>What is this DAG all about?</i>	Call DAG function with arguments of DAG characteristics- name, run schedule, etc
<i>What is each task you want to perform?</i>	<b>Call each <u>operator</u>, and define each task with the arguments</b>
<i>What order should things happen?</i>	Define dependencies (task >> next_task)

For further reading: <https://airflow.apache.org/docs/stable/tutorial.html#it-s-a-dag-definition-file>

# Operators define individual actions for Airflow to complete.

**A task in Airflow is asking for an action to happen. An Operator is a class that you call to initialize the action.**

What kind of task?

- ❖ Run a bash script?
- ❖ Run a python script?
- ❖ Take data from S3 and drop it in Redshift?
- ❖ Run a command on Redshift?
- ❖ Something else?

**Many operators are already defined, and you can use these off the shelf.**

# Operators can be customized or written to meet your needs.

**If you need a task that has no built-in operator:**

- ❖ Start with an existing operator, look at the source code.
- ❖ Can you adapt it? Copy and paste the source, make your edits, save in your environment.
- ❖ If not, use it as a model and write your own from scratch.



# Ask Airflow to Initialize an Athena Query.

Locate within the dag, get all the default args

```
find_records = AWSAthenaOperator(  
    task_id="find_records",  
    dag=dag,  
    query=hdq.query_text.format(datestr="{{ ds }}"),  
    aws_conn_id="aws_default",  
    output_location="s3://my_bucket/my_folder/",  
    database="production",  
)
```

A built-in Airflow macro!

Most arguments are specific to the AWS Athena operator

# Take these steps to make Airflow work for you.

## **Understand Acyclic Graphing (order of tasks)**

- ❖ What really needs to happen first?
- ❖ Think creatively, so you can use the parallelization benefits

## **Use built in credentials management**

- ❖ Let Airflow manage the creds for database connections

## **Practice navigating the GUI**

- ❖ Find your logs, read the graphs to debug easily

# Mistakes I've Made with Airflow

## **Lousy memory management**

- ❖ Holding large data volumes in memory is unnecessary expense - your workers may have issues

## **Too much repetition in DAG code**

- ❖ Think programmatically, be concise
- ❖ Use the default args, try Jinja templating

## **Sloppy Logging/Failure Tracking**

- ❖ Quietly record metadata on results of your jobs in case you need to postmortem
- ❖ Don't page for every little error, but make sure someone is informed if a job fails

# Practical Examples



# Success can be measured concretely.

Build and populate a data warehouse	Run scheduled data analysis tasks and models	Update and manage data warehouse	Make data warehouse easy to use
Scheduled, transparent jobs  Built-in checks for data quality and typing	Train models on a schedule or with one click  Store results in Redshift hands-free	Use Redshift features to minimize resource use  Monitor and log changes to data warehouse	Redshift <> R/Python  Direct querying with SQL from IDE

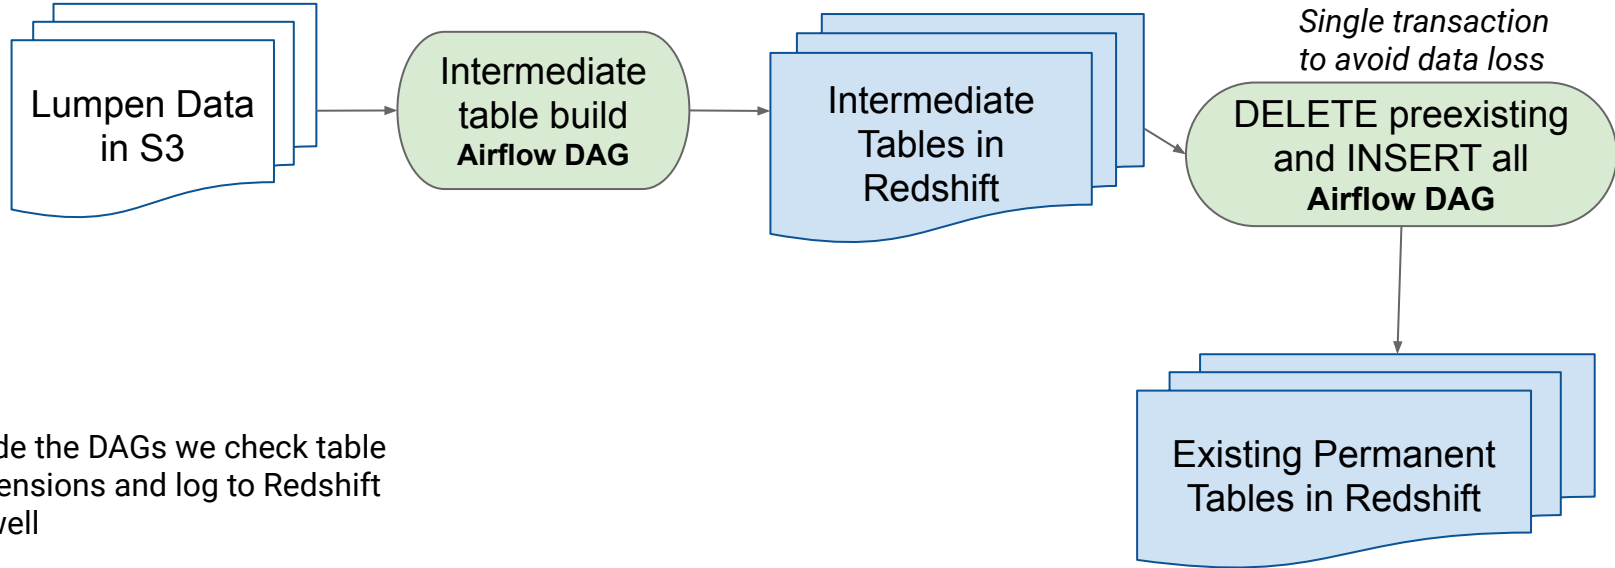
As in any technical implementation, you should have acceptance criteria to define success.

# Example New Data Ingestion Workflow



Airflow is more or less running COPY for you on a schedule with robust error handling.

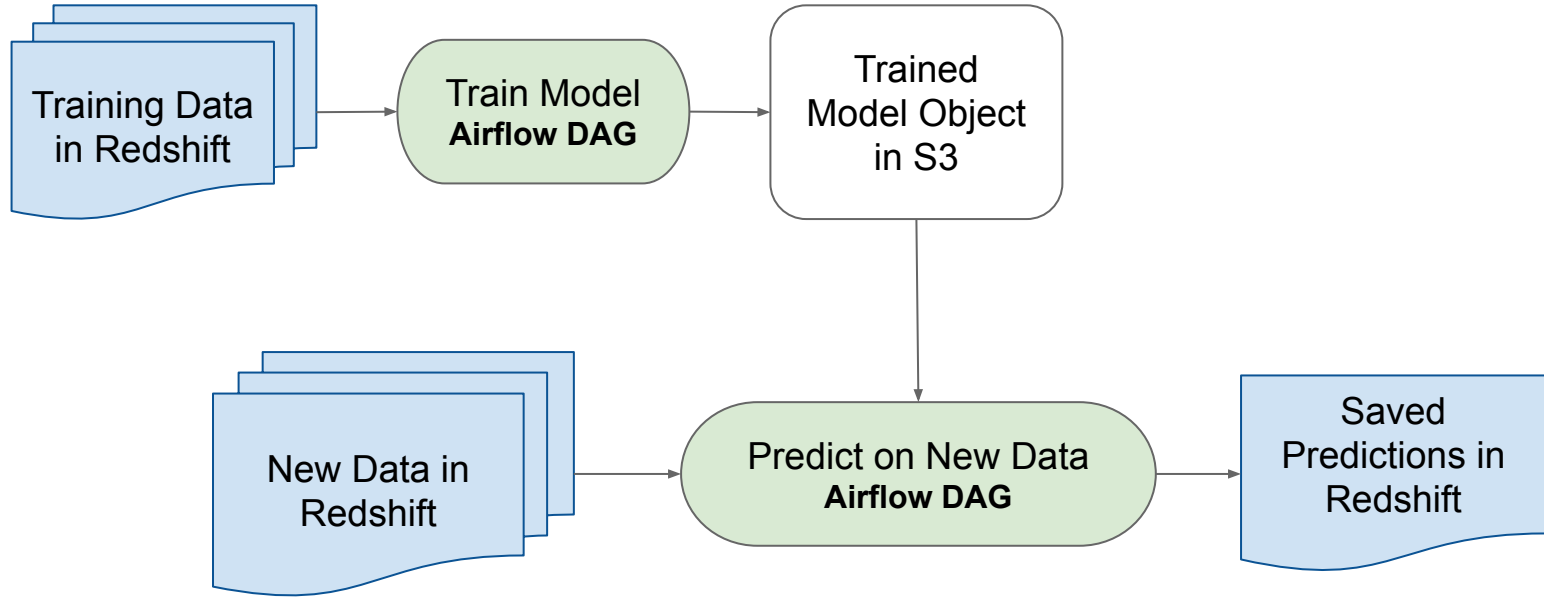
# Example Data Updating Workflow



Inside the DAGs we check table dimensions and log to Redshift as well

Why DELETE + INSERT, not UPDATE? Put a pin in it, I will explain in a moment!

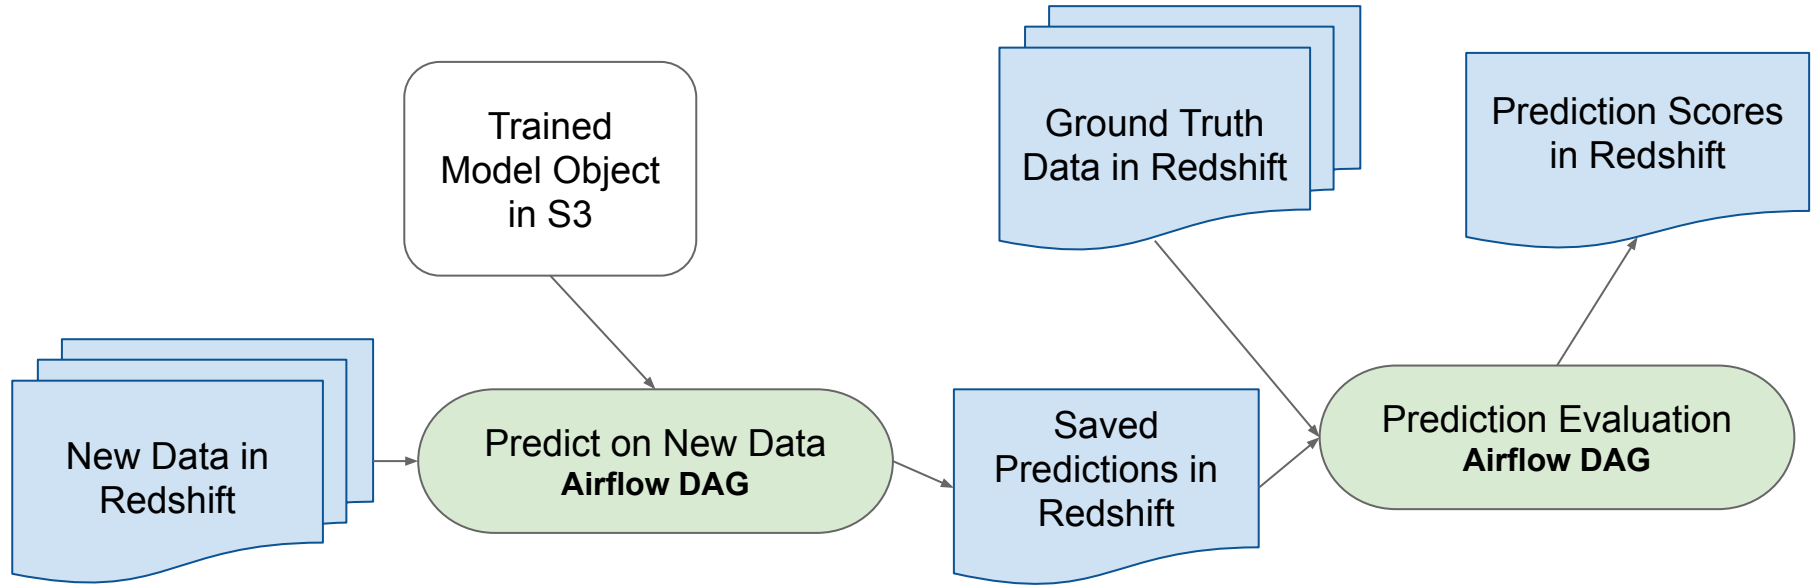
# Example ML Workflow



Airflow runs two different jobs to make the pipeline work – training will be less frequent.



# Example Prediction Evaluation Workflow



This can be one DAG with many tasks in it, depending on how you get ground truth data.

# A Moment on UPDATE

Redshift would rather not

**TL;DR: Redshift has to process every column in order to update one row**

Long answer:

Compression and columnar data storage means that it's unusually resource intensive to update row by row. The resource intensiveness and slowness is dramatic.

**Instead:**

- ❖ Create Temp Table 1 with the new row values
- ❖ Delete any rows from Existing Table that would have been amended
- ❖ INSERT rows from Temp Table 1 into Existing Table
- ❖ Delete Temp Table 1 as cleanup

**INSERT is much easier for Redshift to handle.**

# Finishing touches to make your life much easier.

Develop **logging and monitoring** processes in the DAGs so that you can catch any problems

- ❖ Table size checks
- ❖ Data quality checks
- ❖ DAG failure monitoring/log review

**Teach end users** to write their own DAGs - empower users to generate value themselves.

**Document, document, document** - make sure you have a data dictionary or API guide.

# Additional Links/Resources

## Redshift

- ❖ Using COPY:  
[https://docs.aws.amazon.com/redshift/latest/dg/r\\_COPY.html](https://docs.aws.amazon.com/redshift/latest/dg/r_COPY.html)
- ❖ Analyze compression:  
[https://docs.aws.amazon.com/redshift/latest/dg/r\\_ANALYZE\\_COMPRESSION.html](https://docs.aws.amazon.com/redshift/latest/dg/r_ANALYZE_COMPRESSION.html)
- ❖ Sort keys:  
[https://docs.aws.amazon.com/redshift/latest/dg/c\\_best-practices-sort-key.html](https://docs.aws.amazon.com/redshift/latest/dg/c_best-practices-sort-key.html)
- ❖ Dist keys:  
[https://docs.aws.amazon.com/redshift/latest/dg/c\\_best-practices-best-dist-key.html](https://docs.aws.amazon.com/redshift/latest/dg/c_best-practices-best-dist-key.html)
- ❖ Optimizing queries:  
[https://docs.aws.amazon.com/redshift/latest/dg/c\\_designing-queries-best-practices.html](https://docs.aws.amazon.com/redshift/latest/dg/c_designing-queries-best-practices.html)
- ❖ [https://docs.aws.amazon.com/redshift/latest/dg/r\\_EXPLAIN.html](https://docs.aws.amazon.com/redshift/latest/dg/r_EXPLAIN.html)

# Additional Links/Resources

## Airflow

- ❖ Tutorial:  
<https://airflow.apache.org/docs/stable/tutorial.html>
- ❖ Operators:  
<https://airflow.apache.org/docs/stable/howto/operator/index.html>
- ❖ Built in macros:  
<https://airflow.apache.org/docs/stable/macros-ref.html>
- ❖ The codebase! <https://github.com/apache/airflow>
- ❖ Structuring a DAG:  
<https://airflow.apache.org/docs/stable/tutorial.html#it-s-a-dag-definition-file>

# Thank You!

Questions?

Stephanie Kirmer  
[www.stephaniekirmer.com](http://www.stephaniekirmer.com)  
[https://github.com/skirmer/airflow\\_plus\\_redshift](https://github.com/skirmer/airflow_plus_redshift)  
 @data\_stephanie

```
class RedshiftOperator(BaseOperator):
```

```
    template_fields = ["sql_statement"]
```

```
    @apply_defaults
```

```
    def __init__(self, redshift_conn_id, sql_statement, *args, **kwargs):
```

```
        self.redshift_conn_id = redshift_conn_id
```

```
        self.sql_statement = sql_statement
```

```
        super().__init__(*args, **kwargs)
```

```
    def execute(self, context):
```

```
        self.hook = PostgresHook(postgres_conn_id=self.redshift_conn_id)
```

```
        conn = self.hook.get_conn()
```

```
        cursor = conn.cursor()
```

```
        log.info("Connected to " + self.redshift_conn_id)
```

```
        cursor.execute(self.sql_statement)
```

```
        cursor.close()
```

```
        conn.commit()
```

```
        log.info("Redshift SQL command completed")
```

```
        return True
```

This lets us pass  
Airflow macros inside  
the query body

Bonus: A custom Airflow Operator written by me